

# The Relational Model: Database Definition and Integrity Constraints

## Chapter 3



# Relational (Data) Model

- **Schema** = a description of data in terms of a data model
- **Instance** = a table, with rows (aka tuples, records), and columns (aka fields, attributes) that match the schema
  - # of rows: cardinality
  - # of columns: degree or arity

Students

sid	name	login	age
13	Lisa	lsimp	40
41	Bart	bart	20

Courses

cid	cname	Cr.
E-484	EECS484	4
E-584	EECS584	3

Enrolled

sid	cid	Grade
41	E-484	A-
13	E-584	A+

Cardinality: 2;

Degree: 4



# Relational (Data) Model

- **Schema** = a description of data in terms of a data model
- **Instance** = a table, with rows (aka tuples, records), and columns (aka fields, attributes) that match the schema
  - # of rows: cardinality
  - # of columns: degree or arity

Students

sid	name	login	age
13	Lisa	lsimp	40
41	Bart	bart	20

Cardinality: 2;  
Degree: 4

Courses

cid	cname	Cr.
E-484	EECS484	4
E-584	EECS584	3

Cardinality: 2;  
Degree: 3

Enrolled

sid	cid	Grade
41	E-484	A-
13	E-584	A+



# Relational (Data) Model

- **Schema** = a description of data in terms of a data model
- **Instance** = a table, with rows (aka tuples, records), and columns (aka fields, attributes) that match the schema
  - # of rows: cardinality
  - # of columns: degree or arity

Students

sid	name	login	age
13	Lisa	lsimp	40
41	Bart	bart	20

Cardinality: 2;  
Degree: 4

Courses

cid	cname	Cr.
E-484	EECS484	4
E-584	EECS584	3


Cardinality: 2;  
Degree: 3

Enrolled

sid	cid	Grade
41	E-484	A-
13	E-584	A+

Cardinality: 2;  
Degree: 3

# New Scenario: Olympic Games



- Some history:
  - Inspired by the ancient Olympic Games, which were held in Olympia, **Greece** (8<sup>th</sup> century BC).



# Example:

## Instance of Athlete Relation

AID	Name	Country	Sport
1	Simone Biles	USA	Gymnastics
2	Gabby Thomas	USA	Track
3	Quan Hongchan	China	Diving

What is the schema?

# Example:

## Instance of Athlete Relation

AID	Name	Country	Sport
1	Mary Lou Retton	USA	Gymnastics
2	Jackie Joyner-Kersey	USA	Track
3	Guo Jingjing	China	Diving

What is the schema?

(aid: `integer`, name: `string`,  
country: `string`, sport: `string`)

# Relational Query Languages

- Supports simple, powerful **querying** of data
- Queries written **declaratively**
  - In contrast to **procedural** methods
- DBMS is responsible for **efficient evaluation**
  - System can optimize for efficient query execution, and still ensure that the answer does not change
- SQL is the standard database query language







# Structured Query Language (SQL)

- **CREATE a Table**
  - Integrity Constraints
- INSERT records
- SELECT records
- UPDATE records
- DELETE records
- CREATE View
- UPDATE View

# Create a Table (Relation)

```
CREATE TABLE table_name (  
    field1 TYPE,  
    field2 TYPE,  
    ... ..  
) ;
```

# Try it out in DuckDB, or in Sqlite3



```
CREATE TABLE table_name  
(  
    field1 TYPE,  
    field2 TYPE,  
    ...  
) ;
```

Download SQLite and DuckDB:  
<https://www.sqlite.org/download.html>  
<https://duckdb.org>

- Types: char(20), integer, real, text, blob...  
Creating a database, e.g.,  
sqlite3 olympics.sdb  
duckdb olympics.ddb



# Creating Relations in SQL

- Create the Athlete relation
  - Domain constraint (type) enforced when tuples added or modified
- Create the Olympics relation
- Create the Compete relation

```
CREATE TABLE Athlete
(aid INTEGER,
 name CHAR(30),
 country CHAR(20),
 sport CHAR(20));
```

```
CREATE TABLE Olympics
(oid INTEGER,
 year INTEGER,
 city CHAR(20));
```

```
CREATE TABLE Compete
(aid INTEGER,
 oid INTEGER);
```

# Structured Query Language (SQL)

- **CREATE a Table**
  - Integrity Constraints
- INSERT records
- SELECT records
- UPDATE records
- DELETE records
- CREATE View
- UPDATE View

# Integrity Constraints (ICs)

- Must be true for *every* instance of the database
  - ICs are specified when schema is defined
  - ICs are checked whenever relations are modified
- A *legal* instance of a relation satisfies *\*all\** specified ICs
  - DBMS must not admit illegal instances



# Example Integrity Constraint

- Specify certain attributes as **keys** of a table
- Reference keys of entities in other tables
- No dangling references are allowed in a database when updates occur to any table



# Integrity Constraint:

## Primary and Candidate Keys

- A **key** for a relation  $R$  is a **minimal** set of attributes  $A_1, \dots, A_n$  that must be unique, i.e.,
  - no two tuples in (**any instance of**)  $R$  have the same key attributes  $A_1, \dots, A_n$ .
- One key designated as **primary key**. Others are **candidate keys**
- **Super key**: Only uniqueness requirement, but no minimal requirement. Note: Every key is a superkey.

### Examples:

- $\{\text{aid}\}$  is the primary key in the Athlete relation
- $\{\text{StudentID}\}$  is the primary key for Students relation
- $\{\text{uniqname}\}$  is a candidate key for Students relation






# PRIMARY KEY Constraint

Two ways to specify a primary key constraint:

```
CREATE TABLE Athlete
(aid INTEGER PRIMARY KEY,
 name CHAR(30),
 country CHAR(20),
 sport CHAR(20));
```

```
CREATE TABLE Athlete
(aid INTEGER,
 name CHAR(30),
 country CHAR(20),
 sport CHAR(20),
PRIMARY KEY(aid));
```






# NOT NULL Constraint

- Example: Disallow null values for name attribute:

```
CREATE TABLE Athlete  
(aid INTEGER PRIMARY KEY,  
  name CHAR(30) NOT NULL,  
  country CHAR(20),  
  sport CHAR(20));
```



- NULL** value implies the value is **unknown** or **inapplicable**.  
Semantics:
- Country and sport can be NULL (unspecified)
- But, name must be specified.



# Primary Keys Properties

- SQL Standard:
  - A primary key attribute cannot be null
- Primary key often an integer ID for efficiency

**Design consideration:** Primary keys should be long-term stable, unique, and, ideally, not sensitive. Avoid addresses, SSN, etc.



# Candidate Keys

- Candidate keys specified using **UNIQUE**.
- One of the candidate keys is chosen as the *primary key*.

```
CREATE TABLE Athlete
  (aid INTEGER,
   name CHAR(30) NOT NULL,
   country CHAR(20),
   sport CHAR(20),
   UNIQUE (name, country),
   PRIMARY KEY (aid));
```

- **UNIQUE** attributes can be NULL, unless NOT NULL is specified
  - Only the NON-NULL attribute values need to be **UNIQUE**.

# Foreign Keys in SQL

- Compete table only has relationships between athletes in the Athlete table and the olympics games in the Olympics Table.

```
CREATE TABLE Compete
    (aid INTEGER,    oid INTEGER,
     PRIMARY KEY    (aid, oid)

    ) ;
```

# Foreign Keys in SQL

- Compete table has relationships only between athletes in the Athlete table and the olympics games in the Olympics Table.


```
CREATE TABLE Compete
    (aid INTEGER,    oid INTEGER,
     PRIMARY KEY    (aid, oid),
     FOREIGN KEY (aid) REFERENCES Athlete,
     FOREIGN KEY (oid) REFERENCES Olympics
    ) ;
```

# Foreign Keys in SQL

- Compete table has relationships only between athletes in the Athlete table and the olympics games in the Olympics Table.

```
CREATE TABLE Compete
(aid INTEGER,    oid INTEGER,
  PRIMARY KEY   (aid, oid),
  FOREIGN KEY (aid) REFERENCES Athlete,
  FOREIGN KEY (oid) REFERENCES Olympics
);
```

```
CREATE TABLE Athlete
(aid INTEGER PRIMARY KEY,
name CHAR(30),
country CHAR(20),
sport CHAR(20));
```



Enforcing foreign key constraints implies **referential integrity** (no dangling references) is achieved.

# Sample Database

- Download from course schedule: [athlete\\_create.sql](#), [athlete\\_drop.sql](#)
- You can load it into Sqlite3 as follows:

```
% sqlite3 athlete.db
.read athlete_create.sql
```

Try a SQL query to print a table, e.g.,  
`SELECT * FROM Athlete;`

Some useful Sqlite3 directives

```
.headers on
.mode column
.schema
.help
.read filename
.quit
```
- DuckDB uses same commands as Sqlite3. Download and try it out
- Oracle sqlplus, uses the following to read a file:  
`START athlete_create.sql`



# Enforcing Integrity Constraints (ICs)

- Whenever we modify the database
  - the DBMS must check for violations of ICs
- Enforcing Domain, Primary Key, Unique ICs is straightforward
  - Reject offending UPDATE / INSERT command



# Enforcing Referential Integrity

- If a Compete tuple is **inserted** with no corresponding Athlete aid:
  - Insert operation is REJECTED!

```
CREATE TABLE Compete
  (aid INTEGER,    oid INTEGER,
   PRIMARY KEY    (aid, oid),
   FOREIGN KEY    (aid) REFERENCES Athlete,
   FOREIGN KEY    (oid) REFERENCES Olympics);
```



# Enforcing Referential Integrity

- If a Compete tuple is **inserted** with no corresponding Athlete aid:
  - Insert operation is REJECTED!
- What if an Athlete tuple is **deleted**?

# Referential Integrity enforced in SQL standard on DELETE/UPDATE

- NO ACTION (also called RESTRICT): disallow action (default)
- CASCADE: also delete all referencing tuples
- SET NULL / SET DEFAULT: sets foreign key value of referencing tuple to NULL or a default value

```
CREATE TABLE Compete
  (aid INTEGER,   oid INTEGER,
   PRIMARY KEY   (aid, oid),
   FOREIGN KEY (aid)
  REFERENCES Athlete
  ON DELETE CASCADE
  ON UPDATE SET NULL) ;
```

What happens to the Compete relation if we add a new athlete (with a new ID) to the Athlete table?

# Where do ICs Come From?

- Based on real-world enterprise being modeled
- An IC is a statement about **all** possible instances!
- We can **check** a database instance to see if an IC is **violated**, but we can **NEVER** infer that an IC is true by looking at an instance.
- Key and foreign key ICs are the most common



# Destroying & Altering Relations

- To destroy the relation Olympics.
  - Schema information and tuples are deleted

```
DROP TABLE Olympics
```
- To alter the Athlete schema by adding a new column

```
ALTER TABLE Athlete  
ADD COLUMN age: INTEGER
```
- What do we put in the new field?
  - A **null** value: 'unknown' or 'inapplicable'

# Relational Model: Summary

- A tabular representation of data
- Simple and intuitive
- The most widely used database model by far
- Integrity constraints can be specified by the DBA, based on application semantics. DBMS checks for violations.
  - Two important ICs: primary and foreign keys
  - We always have domain constraints  
e.g. `INTEGER` fields must always contain integer values
- Views can be used for external schemas, and provide logical data independence

Students

sid	name	login	age
13	Lisa	lsimp	40
41	Bart	bart	20

Courses

cid	cname	Cr.
E-484	EECS484	4
E-584	EECS584	3

Enrolled

sid	cid	Grade
41	E-484	A-
13	E-584	A+

