

Logical Database Design: Mapping ER to Relational

Chapter 3, Section 3.5


Learning Objectives

- ER Model used for conceptual design
- Relational Model implemented by modern DBMS
- Learning Objectives: Translate ER diagram to Relational schema

Recall ER Constructs


- Basic Constructs
 - Entity Sets
 - Relationship Sets
 - Attributes (of entities and relationships)
- Additional Constructs
 - ISA Hierarchies
 - Weak Entities
 - Aggregation
- Integrity Constraints
 - Key constraints and Participation constraints
 - ISA hierarchies

Entity Sets to Tables



```
CREATE TABLE Citizens
    (cid    INTEGER,
     name   CHAR(20),
     bday   DATE,
     PRIMARY KEY (cid))
```

Question??




```
CREATE TABLE Citizens  
    (cid    INTEGER  
     name   CHAR(20),  
     bday   DATE,  
     PRIMARY KEY (cid));
```

Can cid have a null value?



- A. Yes
- B. No
- C. Depends


Relationship Sets to Tables



- Foreign key: keys from participating entity sets
- Descriptive attributes

```
CREATE TABLE Votes (  
  cid      INTEGER,  
  iid      INTEGER,  
  when     DATE,  
  PRIMARY KEY (cid, iid),  
  FOREIGN KEY (cid)  
    REFERENCES Citizens,  
  FOREIGN KEY (iid)  
    REFERENCES Initiative);
```

Impact of attribute names

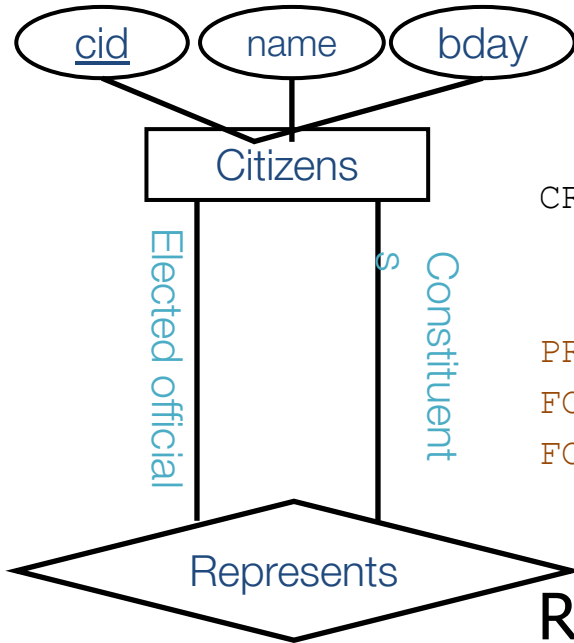


```
CREATE TABLE Votes (
    citizenid    INTEGER,
    iid          INTEGER,
    when         DATE,
    PRIMARY KEY (citizenid, iid),
    FOREIGN KEY (citizenid)
        REFERENCES Citizens(cid),
    FOREIGN KEY (iid)
        REFERENCES Initiative);
```

- Foreign key: keys from participating entity sets
- Descriptive attributes

Note that you need to specify the cid column in REFERENCES.


Relationship Sets to Tables



```
CREATE TABLE Represents(  
    elected_cid  INTEGER,  
    cons_cid    INTEGER,  
    PRIMARY KEY (elected_cid, cons_cid),  
    FOREIGN KEY (elected_cid) REFERENCES Citizens(cid),  
    FOREIGN KEY (cons_cid) REFERENCES Citizens(cid));
```


Represents is a Set of $\{(elected_cid, cons_cid), \dots\}$

Key Constraints



- Notice the arrow from Citizens to Votes

Approach 1: Key Constraint




- Three Tables (Citizens, Votes, Initiative). Votes table changes:

```
CREATE TABLE Votes
(  cid      INTEGER,
   iid      INTEGER,
   when     DATE
   -- Primary key and foreign key constraints?

);
```


Approach 1: Key Constraint



- Three Tables (Citizens, Votes, Initiative). Votes table changes:

```
CREATE TABLE Votes
(  cid      INTEGER NOT NULL,  -- NOT NULL is optional here
  iid      INTEGER NOT NULL,
  when     DATE,
  PRIMARY KEY (cid),  -- implies NOT NULL constraint
  FOREIGN KEY (cid) REFERENCES Citizens,
  FOREIGN KEY (iid) REFERENCES Initiative)
;
```

Approach 2: Key Constraint




- Approach 2: Only two tables: Citizen_Votes and Initiative tables

```
CREATE TABLE Citizens_Votes (  
    cid        INTEGER,  
    name       CHAR(20),  
    bday       DATE,  
    when       DATE,          -- can be NULL  
    iid        INTEGER,       -- can be NULL  
    PRIMARY KEY (cid),  
    FOREIGN KEY (iid) REFERENCES Initiative);
```

Each citizen can only vote once, so OK to fold 'Votes' relationship into 'Citizens' entity and use a single table for both

Question: Which approach better?



- Approach 2: Only two tables: Citizen_Votes and Initiative tables


Approach 1 (separate relationship table)

```
CREATE TABLE Votes
(  cid    INTEGER,
   iid    INTEGER NOT NULL,
   when   DATE,
   PRIMARY KEY (cid),
   FOREIGN KEY (cid) REFERENCES Citizens,
   FOREIGN KEY (iid) REFERENCES Initiative)
);
```

Approach 2 (fold tables)

```
CREATE TABLE Citizens_Votes (
    cid    INTEGER,
    name   CHAR(20),
    bday   DATE,
    when   DATE,
    iid    INTEGER,
    PRIMARY KEY (cid),
    FOREIGN KEY (iid) REFERENCES Initiative)
);
```


Key Constraints



- What about folding everything into **one** table?



Key Constraints



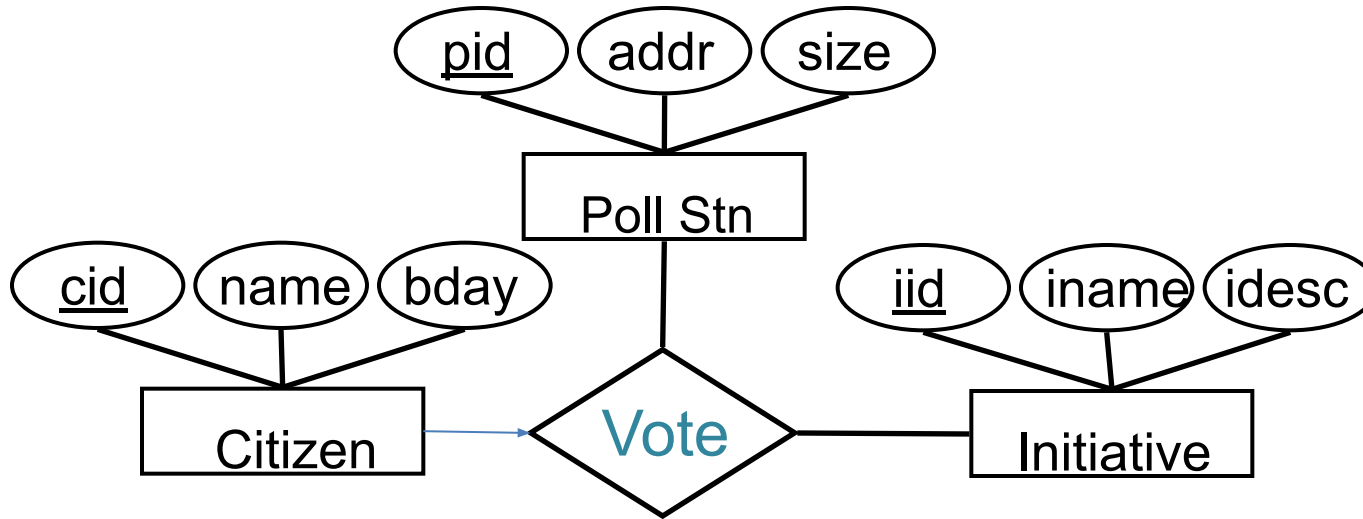
- What about folding everything into **one** table?
- No! This is bad design.



e.g., For every citizen that votes for Initiative A, we have to store iname, idesc information.

(cid, name, bday, when, iid,
iname, idesc) <= REDUNDANCY!

N-ary relationship Example



APPROACH 2:

```
CREATE TABLE Citizens_Votes
```

```
(  cid      INTEGER,  
   name     CHAR(20),  
   bday     DATE,  
   pid      INTEGER,
```

-- NOTE: allowed to be

```
NULL
```

```
   iid      INTEGER,
```

-- NOTE: allowed to be

```
NULL
```


```
PRIMARY KEY (cid),
```

```
FOREIGN KEY (iid) REFERENCES Initiative,
```

How many tables
Using Approach 1?

Using Approach 2?

Participation Constraints




Try Approach 2

```
CREATE TABLE Citizen_Votes (  
    cid      INTEGER,  
    name     CHAR(20),  
    bday     DATE,  
    when     DATE,  
    iid      INTEGER,  
    PRIMARY KEY (cid),  
    FOREIGN KEY (iid) REFERENCES Initiative);
```



What ER constraint is missing?

Participation Constraints




Try Approach 2

```
CREATE TABLE Citizen_Votes(  
    cid      INTEGER,  
    name     CHAR(20),  
    bday     DATE,  
    when     DATE,  
    iid      INTEGER NOT NULL,  
    PRIMARY KEY (cid),  
    FOREIGN KEY (iid) REFERENCES Initiative);
```



What ER constraint is missing?

Participation Constraints




Try Approach 1?



- Citizens table, Votes table, Initiative Table
- How do you guarantee that every citizen in the Citizens table also appears in the Votes table?
- Unfortunately, this is hard to enforce in SQL. Requires multi-table constraints, which are usually not supported in DBMSes.

Mapping Participation Constraints (1-to-1)




```
CREATE TABLE RAB (  
    r1 Integer,  
    a1 Integer,  
    a2 Integer,  
    a3 Integer,  
    b1 Integer,  
    b2 Integer,  
    b3 Integer ...)
```

Key constraints?




Mapping Participation Constraints (1-to-1)



```
CREATE TABLE RAB (  
    r1 Integer,  
    a1 Integer,  
    a2 Integer,  
    a3 Integer,  
    b1 Integer NOT NULL,  
    b2 Integer,  
    b3 Integer,  
    UNIQUE (b1), PRIMARY KEY (a1) )
```

Participation Constraint Only




Try Approach 2?



- Doesn't work since a Citizen can vote multiple times.

Participation Constraint Only




Try Approach 1?



- Citizens table, Votes table, Initiative Table
- How to guarantee that every Citizen appears in the Votes table?
- Unfortunately, this is hard to enforce in SQL. Requires multi-table constraints, which are usually not supported in DBMSes.
- **Typical Practical Solution:** Ignore participation-only constraints – don't bother enforcing them in the database.

Use Approach 1

Weak Entities




- Use Approach 2: Combine weak entity and owning relationship into one relation
 - Delete all weak entities when an owner entity is deleted.

```
CREATE TABLE Expert_Staff (  
    ename      CHAR(20) ,  
    jobDesc    CHAR(40) ,  
    salary     REAL ,  
    cid        INTEGER ,  
    PRIMARY KEY (ename, cid) ,  
    FOREIGN KEY (cid) REFERENCES PR-Cands ON DELETE CASCADE )
```


ISA Hierarchies: General Approach

- Three relations:

- Citizens(ssn, name, bday)
- PR-Cands(ssn, waddr, budget)
- President(ssn, from, to)
- Add foreign key constraints




- Queries:

- Involving all citizens => Easy
- Involving just PR-Cands => may need to **join** PR-cands with Citizens to get the needed attributes

Aggregation – E.g., Monitors

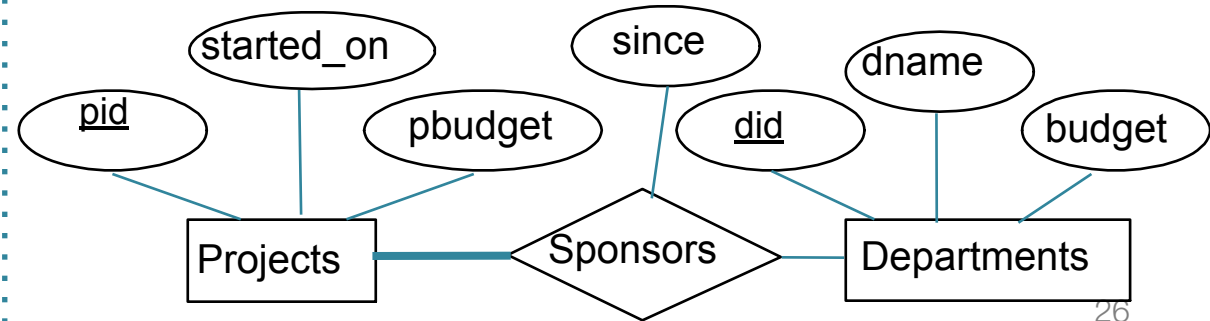


- Tables needed (Approach 1):
 - Projects (pid, started, pbudget)
 - Departments (did, dname, budget)
 - Sponsors (pid, did, since)
 - Manager(mid, name, rank)
 - Monitors(pid, did, mid, until)



Exercise:

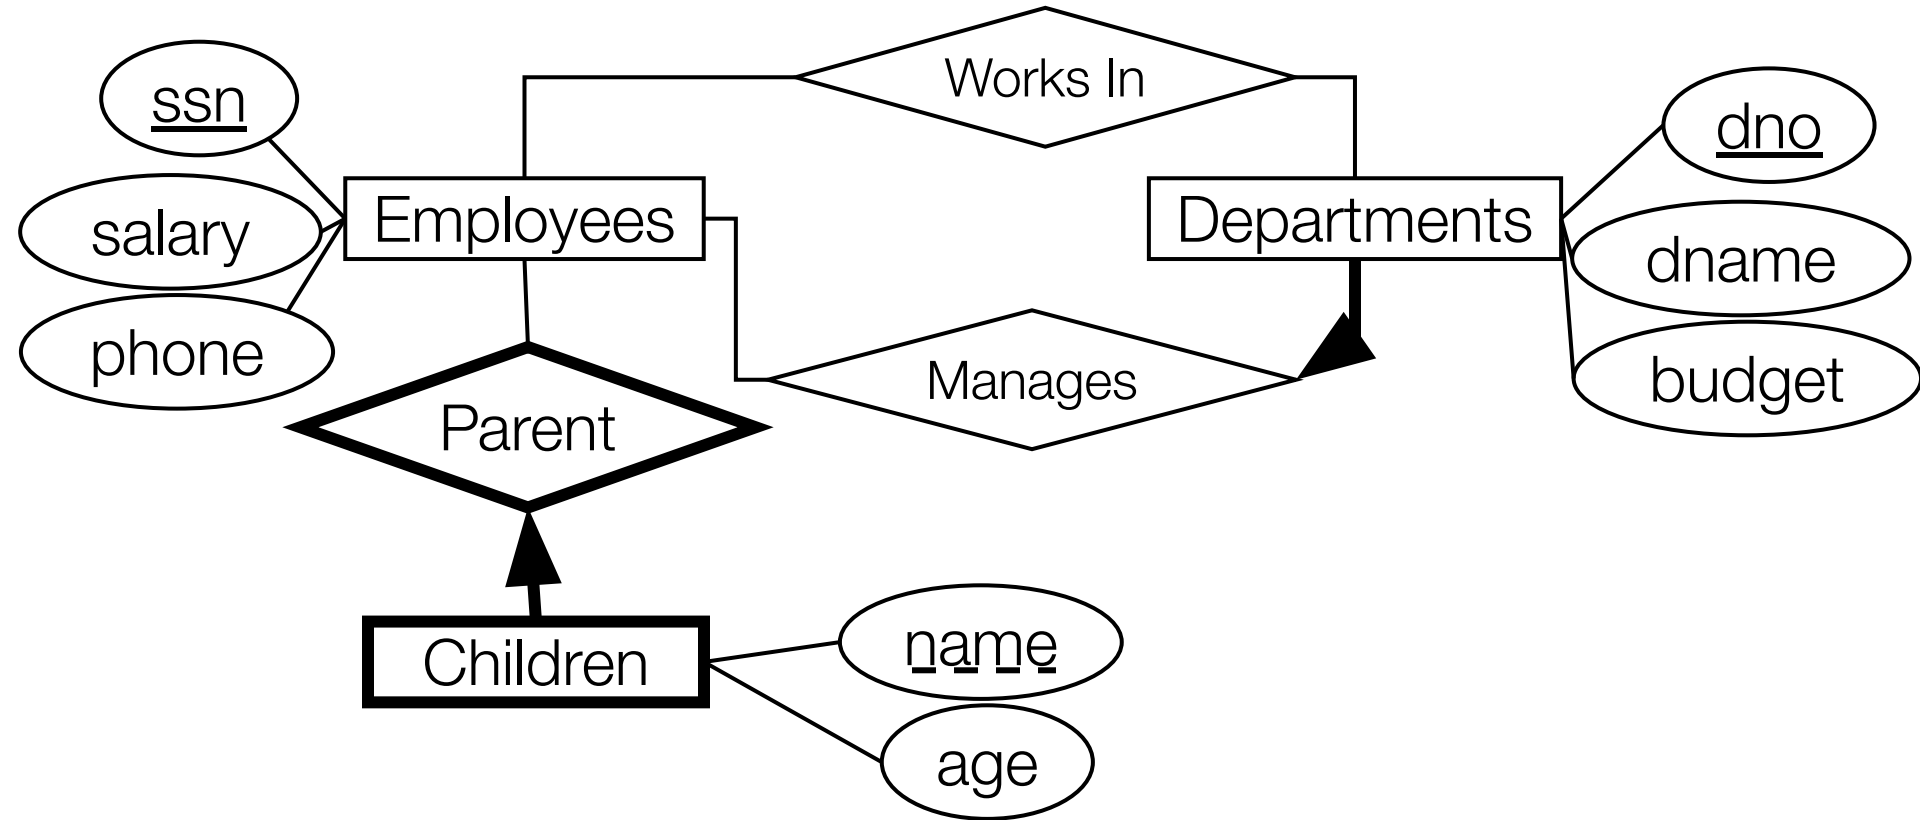
Try Approach 2



Exercise – Part 1

- A **company database** needs to store information about
 - **employees** (identified by ssn, with salary and phone attributes),
 - **departments** (identified by dno, with dname and budget attributes), and
 - **children of employees** (with name and age attributes).
- Employees work in (zero or more) departments
- Each department is managed by exactly one employee
- A child must be identified uniquely by name when the parent (who is an employee; assume only one parent works for the company) is known.
- We are not interested in information about a child once the parent leaves the company.
- Draw an ER diagram that captures this information

ER Diagram (one solution)



Exercise – Part 2

- Write SQL statements to create the corresponding relations, and to capture as many of the constraints as possible.

SQL DDL – One solution



```
CREATE TABLE Employees (  
  ssn INTEGER,  
  salary REAL,  
  phone CHAR(10),  
  PRIMARY KEY(ssn))
```

```
CREATE TABLE Works (  
  ssn INTEGER,  
  dno INTEGER,  
  PRIMARY KEY (ssn, dno),  
  FOREIGN KEY (ssn)  
    REFERENCES employees,  
  FOREIGN KEY (dno)  
    REFERENCES departments)
```

```
CREATE TABLE Departments (  
  dno INTEGER,  
  dname CHAR(20),  
  budget real,  
  manager INTEGER NOT NULL,  
  PRIMARY KEY (dno),  
  FOREIGN KEY (manager)  
    REFERENCES employees(ssn))
```

```
CREATE TABLE Children (  
  name CHAR(20),  
  age REAL,  
  parent INTEGER NOT NULL,  
  PRIMARY KEY(name, parent),  
  FOREIGN KEY(parent)  
    REFERENCES employees(ssn)  
    ON DELETE CASCADE)
```

Integrity Constraints

- Describe conditions that must be satisfied by every legal instance
- Types of integrity constraints
 - Domain constraints
 - Primary key constraints
 - Foreign key constraints
 - General constraints

Table Constraints

- More general than key constraints
- Can use a query to express constraint
 - Constraints checked each time a table is updated
 - CHECK constraint always true for empty relation

```
CREATE TABLE Athlete
(aid INTEGER PRIMARY KEY,
name CHAR(30),
age INTEGER,
country CHAR(20),
sport CHAR(20)),
CHECK ( age >= 18 AND age <= 80 ));
```


Try it out in sqlplus or sqlite or duckDB

```
CREATE TABLE Sailors
( sid      INTEGER,
  sname    CHAR(10),
  rating   INTEGER,
  age      REAL,
  PRIMARY KEY (sid),
  CHECK (rating >= 1 AND rating <= 10));

INSERT INTO Sailors VALUES (1, 's1', 11, 25);

Do you get a constraint violation error?
```

```
CREATE TABLE Sailors
( sid      INTEGER,
  sname    CHAR(10),
  rating   INTEGER,
  age      REAL,
  PRIMARY KEY (sid),
  CHECK (rating >= 1 AND rating <= 10));

INSERT INTO Sailors VALUES (1, 's1', 11, 25);

> Error: CHECK constraint failed: Sailors
```



More general CHECKs

- SQL standard allows cross-table CHECK constraints
- But, they are not supported in most systems – expensive to enforce.
- Thus, we will not worry about them.

Active Databases & Triggers

Trigger: Procedure that starts automatically if specified changes occur to the DBMS

- Three parts:
 - Event (activates the trigger)
 - Condition (test that is run when the trigger is activated)
 - Action (what happens if the trigger runs)
 - Before and After Triggers
- Trigger Execution
 - Row-level Triggers: Once per modified row
 - Statement-level Triggers: Once per SQL statement

Example – Log entries

```
CREATE TABLE Employee (  
  Empno NUMBER PRIMARY KEY,  
  Ename VARCHAR2(10),  
  Job VARCHAR2(9),  
  Mgr NUMBER(4),  
  Hiredate DATE,  
  Sal NUMBER(7,2));
```

```
CREATE TABLE Emp_log (  
  Emp_id NUMBER,  
  Log_date DATE,  
  New_salary NUMBER,  
  Log_entry VARCHAR2(20));
```

```
CREATE OR REPLACE TRIGGER Log_salary_increase AFTER UPDATE  
ON Employee FOR EACH ROW WHEN (new.Sal > 1000)  
BEGIN  
  INSERT INTO Emp_log (Emp_id, Log_date, New_salary, Log_entry)  
    VALUES (:new.Empno, SYSDATE, :new.SAL, 'NEW SAL');  
END;
```

This is a row-level trigger (triggered per row update). If new salary above 1000, log entry is added.

E.g. by executing this SQL: `UPDATE Emp_tab SET Sal = Sal + 1000.0 WHERE Job = 'Professor';`

Example – Log entries

```
CREATE TABLE Employee (  
  Empno NUMBER PRIMARY KEY,  
  Ename VARCHAR2(10),  
  Job VARCHAR2(9),  
  Mgr NUMBER(4),  
  Hiredate DATE,  
  Sal NUMBER(7,2));
```

```
CREATE TABLE Emp_log (  
  Log_date DATE,  
  Log_entry VARCHAR2(20));
```

```
CREATE OR REPLACE TRIGGER Log_emp_update  
  AFTER UPDATE ON Employee  
BEGIN  
  INSERT INTO Emp_log (Log_date, Log_entry)  
    VALUES (SYSDATE, 'Employee Table CHANGED');  
END;
```

This is a statement-level trigger. The trigger will fire once per update statement.

Triggers


- When condition is checked, success/failure can be used to trigger arbitrary actions.
- Used for many things:
 - Check complex actions (such as credit limit in a shopping application)
 - Generate logs for auditing and security checks.

Recall CASCADE constraints

```
CREATE TABLE Athlete  
(aid INTEGER PRIMARY KEY,  
 name CHAR(30),  
 country CHAR(20),  
 sport CHAR(20));
```

```
CREATE TABLE Olympics  
(oid INTEGER PRIMARY KEY,  
 year INTEGER,  
 city CHAR(20));
```

```
CREATE TABLE Compete  
(aid INTEGER,  
 oid INTEGER,  
 PRIMARY KEY (aid, oid),  
 FOREIGN KEY (aid) REFERENCES Athlete  
     ON DELETE CASCADE,  
 FOREIGN KEY (oid) REFERENCES Olympics  
 );
```



CASCADE Using Triggers

```
CREATE TABLE Compete
  (aid INTEGER,
   oid INTEGER,
   PRIMARY KEY  (aid, oid),
   FOREIGN KEY  (aid) REFERENCES Athlete,
   FOREIGN KEY  (oid) REFERENCES Olympics);
```

```
CREATE OR REPLACE TRIGGER cascade_on_delete
AFTER  DELETE ON Athlete
FOR EACH ROW
BEGIN
  DELETE FROM Compete
  WHERE Compete.aid = :OLD.aid;
END;
```



Trying out triggers



- Try out the file [athlete_trigger_cascade.sql](#) in sqlplus
- Also try it out in sqlite to see if it supports triggers the same way
- Try removing a row from athlete. Does it cascade to Compete via the trigger?
- Try dropping the trigger:
 - `DROP TRIGGER triggername;`
- What happens now on deleting a row from Athlete?

Triggers: Pitfalls and Pain

- Triggers can be recursive!
 - Chain of triggers can be hard to predict, which makes triggers difficult to understand and debug
- Errors with “mutating” table
 - A table that is currently being modified by an UPDATE, DELETE, or INSERT statement, or a table that might be updated by the effects of a DELETE CASCADE constraint
 - The session that issued the triggering statement cannot query or modify a mutating table
 - Used to prevent a trigger from seeing inconsistent data